

# Big Data



## Big Data Engineering with Hadoop & Spark

Kafka Introduction



## Session 22: Assignment 22.1

---

This assignment is aimed at consolidating the concepts that was learnt during the Apache Kafka session of the course.

## Initial Definitions:

**Topic:** A topic is a category or feed name to which records are published. Topics in Kafka are always multi-subscriber; that is, a topic can have zero, one, or many consumers that subscribe to the data written to it.

**Producers:** Producers publish data to the topics of their choice. The producer is responsible for choosing which record to assign to which partition within the topic. This can be done in a round-robin fashion simply to balance load or it can be done according to some semantic partition function.

**Consumers:** Consumers label themselves with a consumer group name, and each record published to a topic is delivered to one consumer instance within each subscribing consumer group.

- If all the consumer instances have the same consumer group, then the records will effectively be load balanced over the consumer instances.
- If all the consumer instances have different consumer groups, then each record will be broadcast to all the consumer processes.

**Zookeeper:** ZooKeeper is a service to which clients can connect to. It provides access to clients to a tree like structure or a hierarchical namespace.

# Task 1:

Create a Kafka topic named KeyLessTopic.

- Start Zookeeper

```
$ $KAFKA_HOME/bin/zookeeper-server-start.sh
  $KAFKA_HOME/config/zookeeper.properties
```

- Start Kafka Broker

```
$ $KAFKA_HOME/bin/kafka-server-start.sh
  $KAFKA_HOME/config/server.properties
```

- Create a topic named KeyLessTopic

```
$ $KAFKA_HOME/bin/kafka-topics.sh --create --topic KeyLessTopic --
  zookeeper localhost:2181 --partitions 1 --replication-factor 1
```

- Inside KeyLessTopic insert following data

```
$ $KAFKA_HOME/bin/kafka-console-producer.sh --broker-list
  localhost:9092 --topic KeyLessTopic
```

- Insert the Data below

```
{"name":"John", "exp":16}
{"name":"Finn", "exp":20}
{"name":"Cylin", "exp":18}
{"name":"Mark", "exp":2}
{"name":"Akshay", "exp":14}
```

```
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-topics.sh --create --topic KeyLessTopic --zookeeper localhost:2181 --partitions 1 --replication-factor 1
Created topic "KeyLessTopic".
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic KeyLessTopic
{"name":"John", "exp":16}
{"name":"Finn", "exp":20}
{"name":"Cylin", "exp":18}
{"name":"Mark", "exp":2}
{"name":"Akshay", "exp":14}
```

# Task 2:

Create a console consumer that reads KeyLessTopic from beginning

- Running a console consumer to read the type KeyLessTopic sent by the producer.

```
$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic
  KeyLessTopic --from-beginning --zookeeper localhost:2181 --
  property print.key=false
```

```
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic KeyLessTopic --from-beginning --zookeeper localhost:2181 --property print.key=false
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap-server] instead of [zookeeper].
{"name":"John", "exp":16}
{"name":"Finn", "exp":20}
{"name":"Cylin", "exp":18}
{"name":"Mark", "exp":2}
{"name":"Akshay", "exp":14}
```

## Task 3:

Create a Kafka topic named KeyedTopic. Inside KeyedTopic insert following data:

- The part before comma(,) should be treated as key and after comma(,) should be treated as value
 

```

{"name":"John"},{"exp":16}
{"name":"Finn"},{"exp":20}
{"name":"Cylin"},{"exp":18}
{"name":"Mark"},{"exp":2}
{"name":"Akshay"},{"exp":14}

```
- Creating Topic
 

```

$ $KAFKA_HOME/bin/kafka-topics.sh --create --topic KeyedTopic --zookeeper localhost:2181 --partitions 1 --replication-factor 1

```
- Creating Producer
 

```

$ $KAFKA_HOME/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic KeyedTopic --property parse.key=true --property key.separator=","

```
- Now, inserting the above-mentioned data

```

[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-topics.sh --create --topic KeyedTopic --zookeeper localhost:2181 --partitions 1 --replication-factor 1
Created topic "KeyedTopic".
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic KeyedTopic --property parse.key=true --property key.separator=","
{"name":"John"},{"exp":16}
{"name":"Finn"},{"exp":20}
{"name":"Cylin"},{"exp":18}
{"name":"Mark"},{"exp":2}
{"name":"Akshay"},{"exp":14}

```

## Task 4:

Create a console consumer that reads KeyedTopic from beginning. The key and value should be separated by '-'

- Creating Consumer
 

```

$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic KeyedTopic -from-beginning --zookeeper localhost:2181 --property print.key=true --property key.separator="-"

```

```

[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic KeyedTopic --from-beginning --zookeeper localhost:2181 --property print.key=true --property key.separator="-"
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap-server] instead of [zookeeper].
{"name":"John"}-{"exp":16}
{"name":"Finn"}-{"exp":20}
{"name":"Cylin"}-{"exp":18}
{"name":"Mark"}-{"exp":2}
{"name":"Akshay"}-{"exp":14}

```